

```
In [1]: import warnings
warnings.filterwarnings('ignore')
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_val_score, KFold,
from sklearn.metrics import accuracy_score, classification_report, confusion_m
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import GradientBoostingClassifier
```

```
In [2]: df = pd.read_csv("creditcard.csv")
```

```
In [3]: df.head(10)
```

Out[3]:

| | Time | V1 | V2 | V3 | V4 | V5 | V6 | V7 | V8 |
|---|------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | 0.0 | -1.359807 | -0.072781 | 2.536347 | 1.378155 | -0.338321 | 0.462388 | 0.239599 | 0.098698 |
| 1 | 0.0 | 1.191857 | 0.266151 | 0.166480 | 0.448154 | 0.060018 | -0.082361 | -0.078803 | 0.085102 |
| 2 | 1.0 | -1.358354 | -1.340163 | 1.773209 | 0.379780 | -0.503198 | 1.800499 | 0.791461 | 0.247676 |
| 3 | 1.0 | -0.966272 | -0.185226 | 1.792993 | -0.863291 | -0.010309 | 1.247203 | 0.237609 | 0.377436 |
| 4 | 2.0 | -1.158233 | 0.877737 | 1.548718 | 0.403034 | -0.407193 | 0.095921 | 0.592941 | -0.270533 |
| 5 | 2.0 | -0.425966 | 0.960523 | 1.141109 | -0.168252 | 0.420987 | -0.029728 | 0.476201 | 0.260314 |
| 6 | 4.0 | 1.229658 | 0.141004 | 0.045371 | 1.202613 | 0.191881 | 0.272708 | -0.005159 | 0.081213 |
| 7 | 7.0 | -0.644269 | 1.417964 | 1.074380 | -0.492199 | 0.948934 | 0.428118 | 1.120631 | -3.807864 |
| 8 | 7.0 | -0.894286 | 0.286157 | -0.113192 | -0.271526 | 2.669599 | 3.721818 | 0.370145 | 0.851084 |
| 9 | 9.0 | -0.338262 | 1.119593 | 1.044367 | -0.222187 | 0.499361 | -0.246761 | 0.651583 | 0.069539 |

10 rows × 31 columns

```
In [4]: df.shape
```

Out[4]: (284807, 31)

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [6]: df.dtypes
```

```
Out[6]: Time      float64  
V1      float64  
V2      float64  
V3      float64  
V4      float64  
V5      float64  
V6      float64  
V7      float64  
V8      float64  
V9      float64  
V10     float64  
V11     float64  
V12     float64  
V13     float64  
V14     float64  
V15     float64  
V16     float64  
V17     float64  
V18     float64  
V19     float64  
V20     float64  
V21     float64  
V22     float64  
V23     float64  
V24     float64  
V25     float64  
V26     float64  
V27     float64  
V28     float64  
Amount   float64  
Class      int64  
dtype: object
```

```
In [7]: df.isna().sum().sort_values()
```

```
Out[7]: Time      0
V28      0
V27      0
V26      0
V25      0
V24      0
V23      0
V22      0
V21      0
V20      0
V19      0
V18      0
V17      0
V16      0
Amount    0
V15      0
V13      0
V12      0
V11      0
V10      0
V9       0
V8       0
V7       0
V6       0
V5       0
V4       0
V3       0
V2       0
V1       0
V14      0
Class     0
dtype: int64
```

```
In [9]: df.describe()
```

```
Out[9]:
```

| | Time | V1 | V2 | V3 | V4 | V5 |
|--------------|---------------|---------------|---------------|---------------|---------------|---------------|
| count | 284807.000000 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 |
| mean | 94813.859575 | 3.918649e-15 | 5.682686e-16 | -8.761736e-15 | 2.811118e-15 | -1.552103e-15 |
| std | 47488.145955 | 1.958696e+00 | 1.651309e+00 | 1.516255e+00 | 1.415869e+00 | 1.380247e+00 |
| min | 0.000000 | -5.640751e+01 | -7.271573e+01 | -4.832559e+01 | -5.683171e+00 | -1.137433e+02 |
| 25% | 54201.500000 | -9.203734e-01 | -5.985499e-01 | -8.903648e-01 | -8.486401e-01 | -6.915971e-01 |
| 50% | 84692.000000 | 1.810880e-02 | 6.548556e-02 | 1.798463e-01 | -1.984653e-02 | -5.433583e-02 |
| 75% | 139320.500000 | 1.315642e+00 | 8.037239e-01 | 1.027196e+00 | 7.433413e-01 | 6.119264e-01 |
| max | 172792.000000 | 2.454930e+00 | 2.205773e+01 | 9.382558e+00 | 1.687534e+01 | 3.480167e+01 |

8 rows × 31 columns



In [10]: `df.describe().T`

Out[10]:

| | count | mean | std | min | 25% | 50% | |
|---------------|----------|---------------|--------------|-------------|--------------|--------------|--------|
| Time | 284807.0 | 9.481386e+04 | 47488.145955 | 0.000000 | 54201.500000 | 84692.000000 | 139320 |
| V1 | 284807.0 | 3.918649e-15 | 1.958696 | -56.407510 | -0.920373 | 0.018109 | 1 |
| V2 | 284807.0 | 5.682686e-16 | 1.651309 | -72.715728 | -0.598550 | 0.065486 | 0 |
| V3 | 284807.0 | -8.761736e-15 | 1.516255 | -48.325589 | -0.890365 | 0.179846 | 1 |
| V4 | 284807.0 | 2.811118e-15 | 1.415869 | -5.683171 | -0.848640 | -0.019847 | 0 |
| V5 | 284807.0 | -1.552103e-15 | 1.380247 | -113.743307 | -0.691597 | -0.054336 | 0 |
| V6 | 284807.0 | 2.040130e-15 | 1.332271 | -26.160506 | -0.768296 | -0.274187 | 0 |
| V7 | 284807.0 | -1.698953e-15 | 1.237094 | -43.557242 | -0.554076 | 0.040103 | 0 |
| V8 | 284807.0 | -1.893285e-16 | 1.194353 | -73.216718 | -0.208630 | 0.022358 | 0 |
| V9 | 284807.0 | -3.147640e-15 | 1.098632 | -13.434066 | -0.643098 | -0.051429 | 0 |
| V10 | 284807.0 | 1.772925e-15 | 1.088850 | -24.588262 | -0.535426 | -0.092917 | 0 |
| V11 | 284807.0 | 9.289524e-16 | 1.020713 | -4.797473 | -0.762494 | -0.032757 | 0 |
| V12 | 284807.0 | -1.803266e-15 | 0.999201 | -18.683715 | -0.405571 | 0.140033 | 0 |
| V13 | 284807.0 | 1.674888e-15 | 0.995274 | -5.791881 | -0.648539 | -0.013568 | 0 |
| V14 | 284807.0 | 1.475621e-15 | 0.958596 | -19.214325 | -0.425574 | 0.050601 | 0 |
| V15 | 284807.0 | 3.501098e-15 | 0.915316 | -4.498945 | -0.582884 | 0.048072 | 0 |
| V16 | 284807.0 | 1.392460e-15 | 0.876253 | -14.129855 | -0.468037 | 0.066413 | 0 |
| V17 | 284807.0 | -7.466538e-16 | 0.849337 | -25.162799 | -0.483748 | -0.065676 | 0 |
| V18 | 284807.0 | 4.258754e-16 | 0.838176 | -9.498746 | -0.498850 | -0.003636 | 0 |
| V19 | 284807.0 | 9.019919e-16 | 0.814041 | -7.213527 | -0.456299 | 0.003735 | 0 |
| V20 | 284807.0 | 5.126845e-16 | 0.770925 | -54.497720 | -0.211721 | -0.062481 | 0 |
| V21 | 284807.0 | 1.473120e-16 | 0.734524 | -34.830382 | -0.228395 | -0.029450 | 0 |
| V22 | 284807.0 | 8.042109e-16 | 0.725702 | -10.933144 | -0.542350 | 0.006782 | 0 |
| V23 | 284807.0 | 5.282512e-16 | 0.624460 | -44.807735 | -0.161846 | -0.011193 | 0 |
| V24 | 284807.0 | 4.456271e-15 | 0.605647 | -2.836627 | -0.354586 | 0.040976 | 0 |
| V25 | 284807.0 | 1.426896e-15 | 0.521278 | -10.295397 | -0.317145 | 0.016594 | 0 |
| V26 | 284807.0 | 1.701640e-15 | 0.482227 | -2.604551 | -0.326984 | -0.052139 | 0 |
| V27 | 284807.0 | -3.662252e-16 | 0.403632 | -22.565679 | -0.070840 | 0.001342 | 0 |
| V28 | 284807.0 | -1.217809e-16 | 0.330083 | -15.430084 | -0.052960 | 0.011244 | 0 |
| Amount | 284807.0 | 8.834962e+01 | 250.120109 | 0.000000 | 5.600000 | 22.000000 | 77 |

| | count | mean | std | min | 25% | 50% | |
|--------------|----------|--------------|----------|----------|----------|----------|---|
| Class | 284807.0 | 1.727486e-03 | 0.041527 | 0.000000 | 0.000000 | 0.000000 | 0 |

```
In [11]: f = df[df['Class'] == 1]
v = df[df['Class'] == 0]
print("There are {} farudlent transactions".format(f.shape[0]))
print('There are {} valid transactions'.format(v.shape[0]))
```

There are 492 farudlent transactions
There are 284315 valid transactions

```
In [12]: f.Amount.describe()
```


```
Out[12]: count      492.000000
mean       122.211321
std        256.683288
min         0.000000
25%         1.000000
50%         9.250000
75%        105.890000
max        2125.870000
Name: Amount, dtype: float64
```

```
In [13]: v.Amount.describe()
```

```
Out[13]: count      284315.000000
mean         88.291022
std         250.105092
min          0.000000
25%          5.650000
50%         22.000000
75%         77.050000
max        25691.160000
Name: Amount, dtype: float64
```


```
In [14]: sns.barplot(data = df, x = 'Class', y = 'Amount')  
plt.ylabel("Mean Transaction Amount")
```

```
Out[14]: Text(0, 0.5, 'Mean Transaction Amount')
```




```
In [15]: corr = df.corr()
fig = plt.figure(figsize = (12,10))
sns.heatmap(corr, cmap = 'inferno')
```

Out[15]: <AxesSubplot:>



```
In [16]: scaler = StandardScaler()
df['Amount'] = scaler.fit_transform(df[['Amount']])
```

```
In [17]: X = df.drop('Class', axis = 1).values
y = df['Class'].values
```

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random
```

```
In [ ]: models = {'KNN': KNeighborsClassifier(),
                 'LogReg': LogisticRegression(),
                 'dtree': DecisionTreeClassifier(),
                 'rforest': RandomForestClassifier(),
                 'GB': GradientBoostingClassifier()}
names = []
scores = []
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    score = accuracy_score(y_pred, y_test)
    names.append(name)
    scores.append(score)
```

```
In [23]: results = pd.DataFrame({'model': names, 'score': scores})
results.sort_values(by='score', ascending=False)
```

Out[23]:

| | model | score |
|---|---------|----------|
| 0 | rforest | 0.999579 |